Attention is all you need

Author: Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I..

Publish: *Advances in Neural Information Processing Systems*
Pp: 5998 - 6008.

Presenter: WENWEI KANG

https://arxiv.org/abs/1706.03762

- Introduction
- Encoder - Decoder
- Transformer
- Evaluation

Introduction

Neural Machine Translation(NMT):

- **Statistical based:** Phrase-based + large LM (Moses)
- **NN based:** Encoder - Decoder (Seq2seq, ConvS2S, ensemble …)

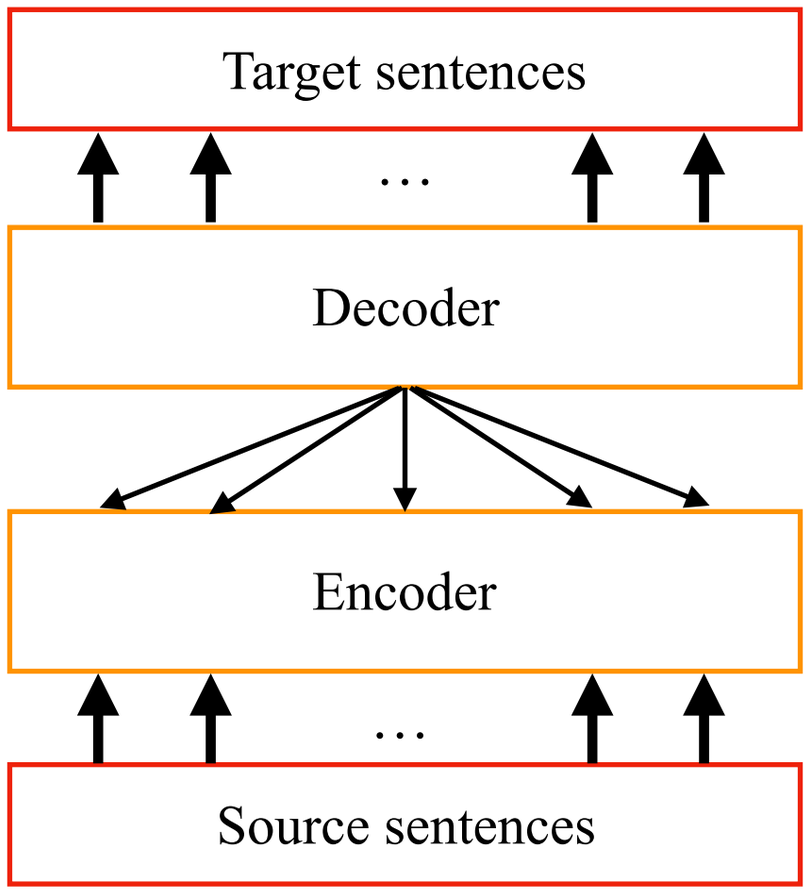

Encoder - Decoder

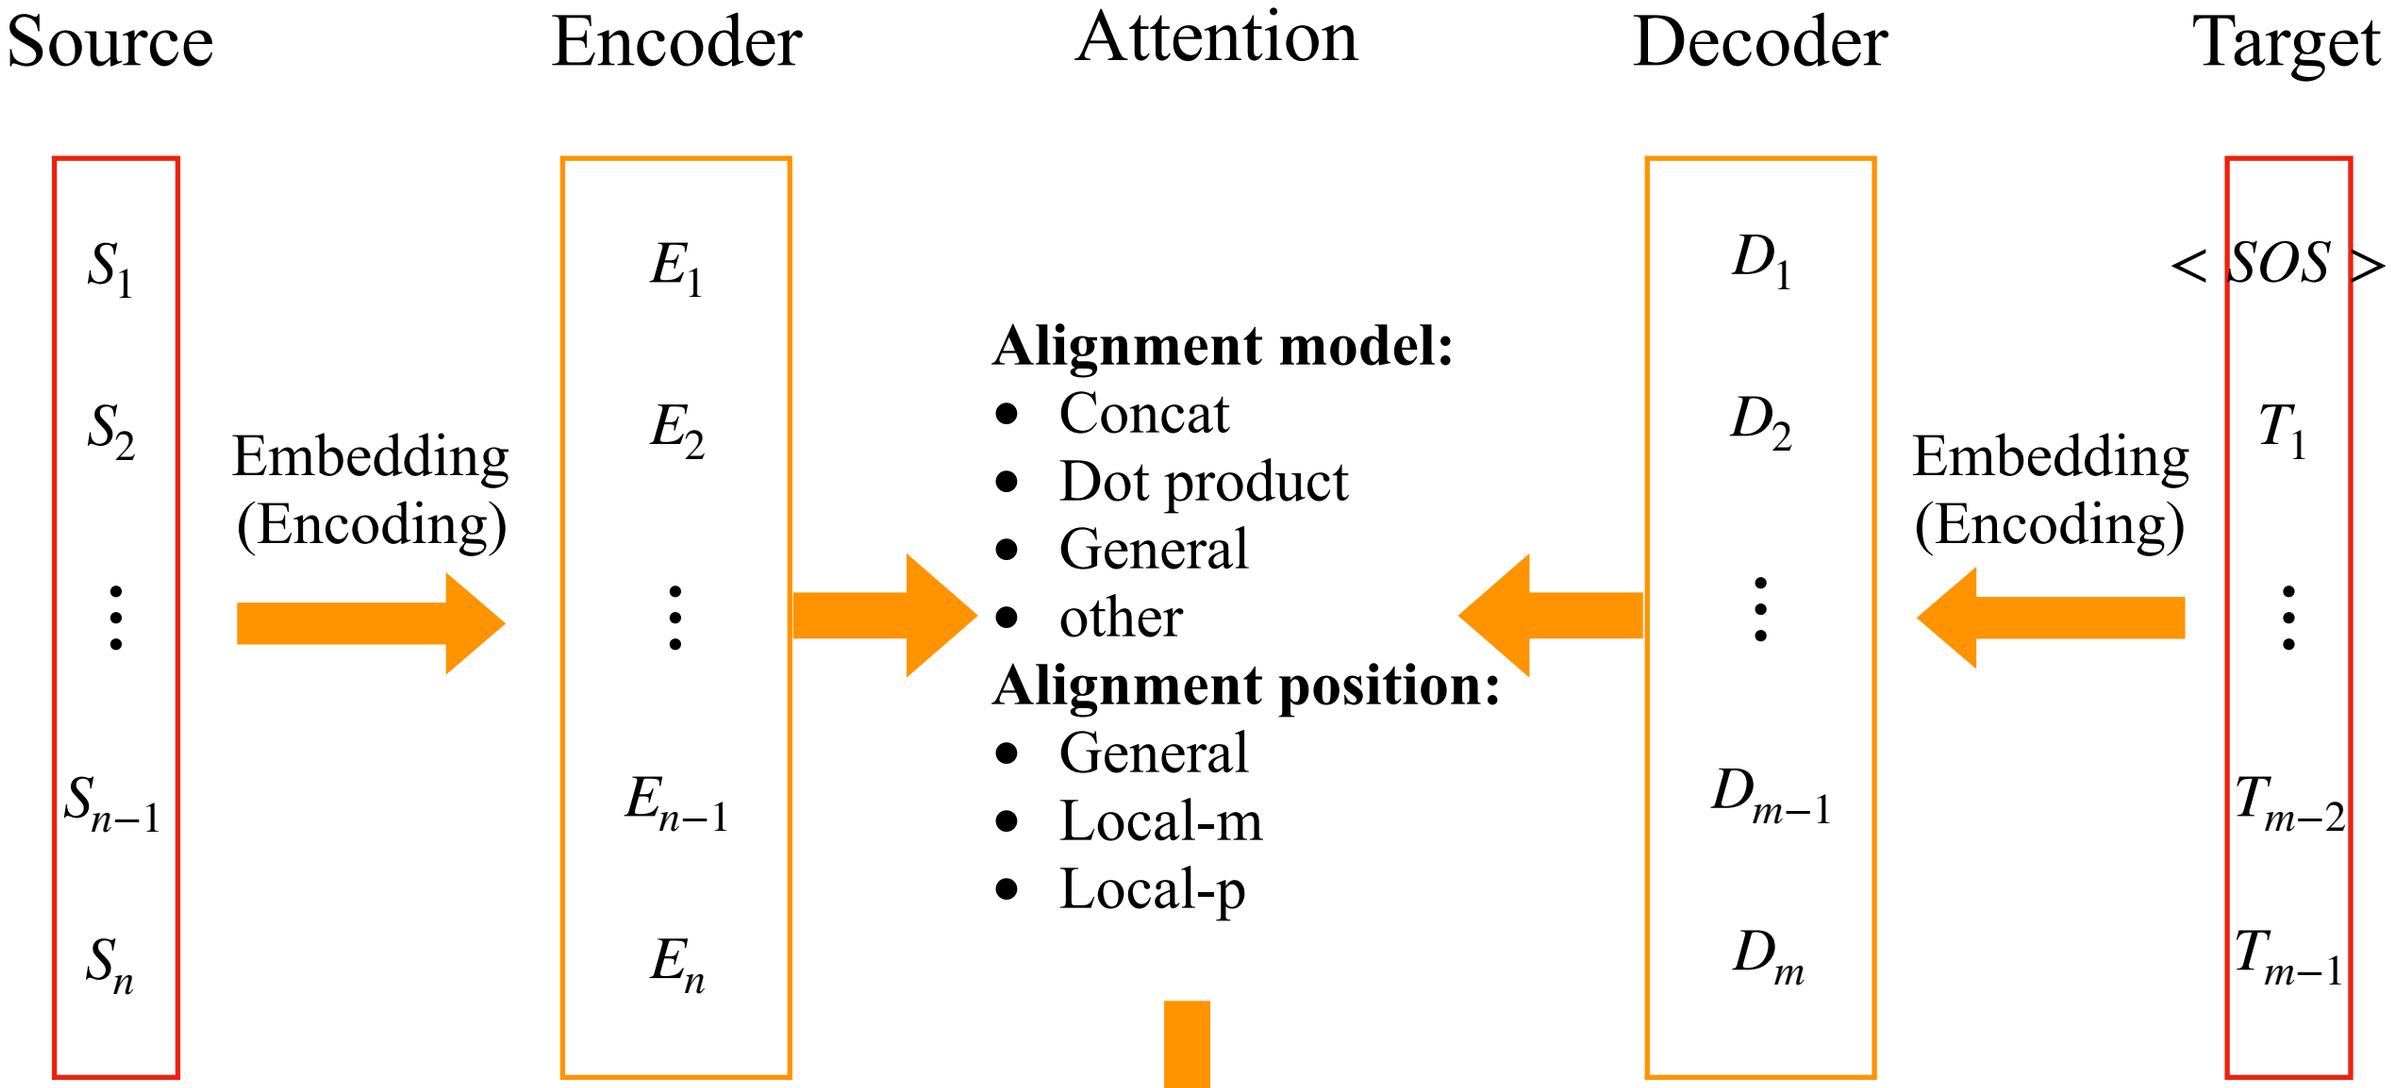

Encoder - Decoder

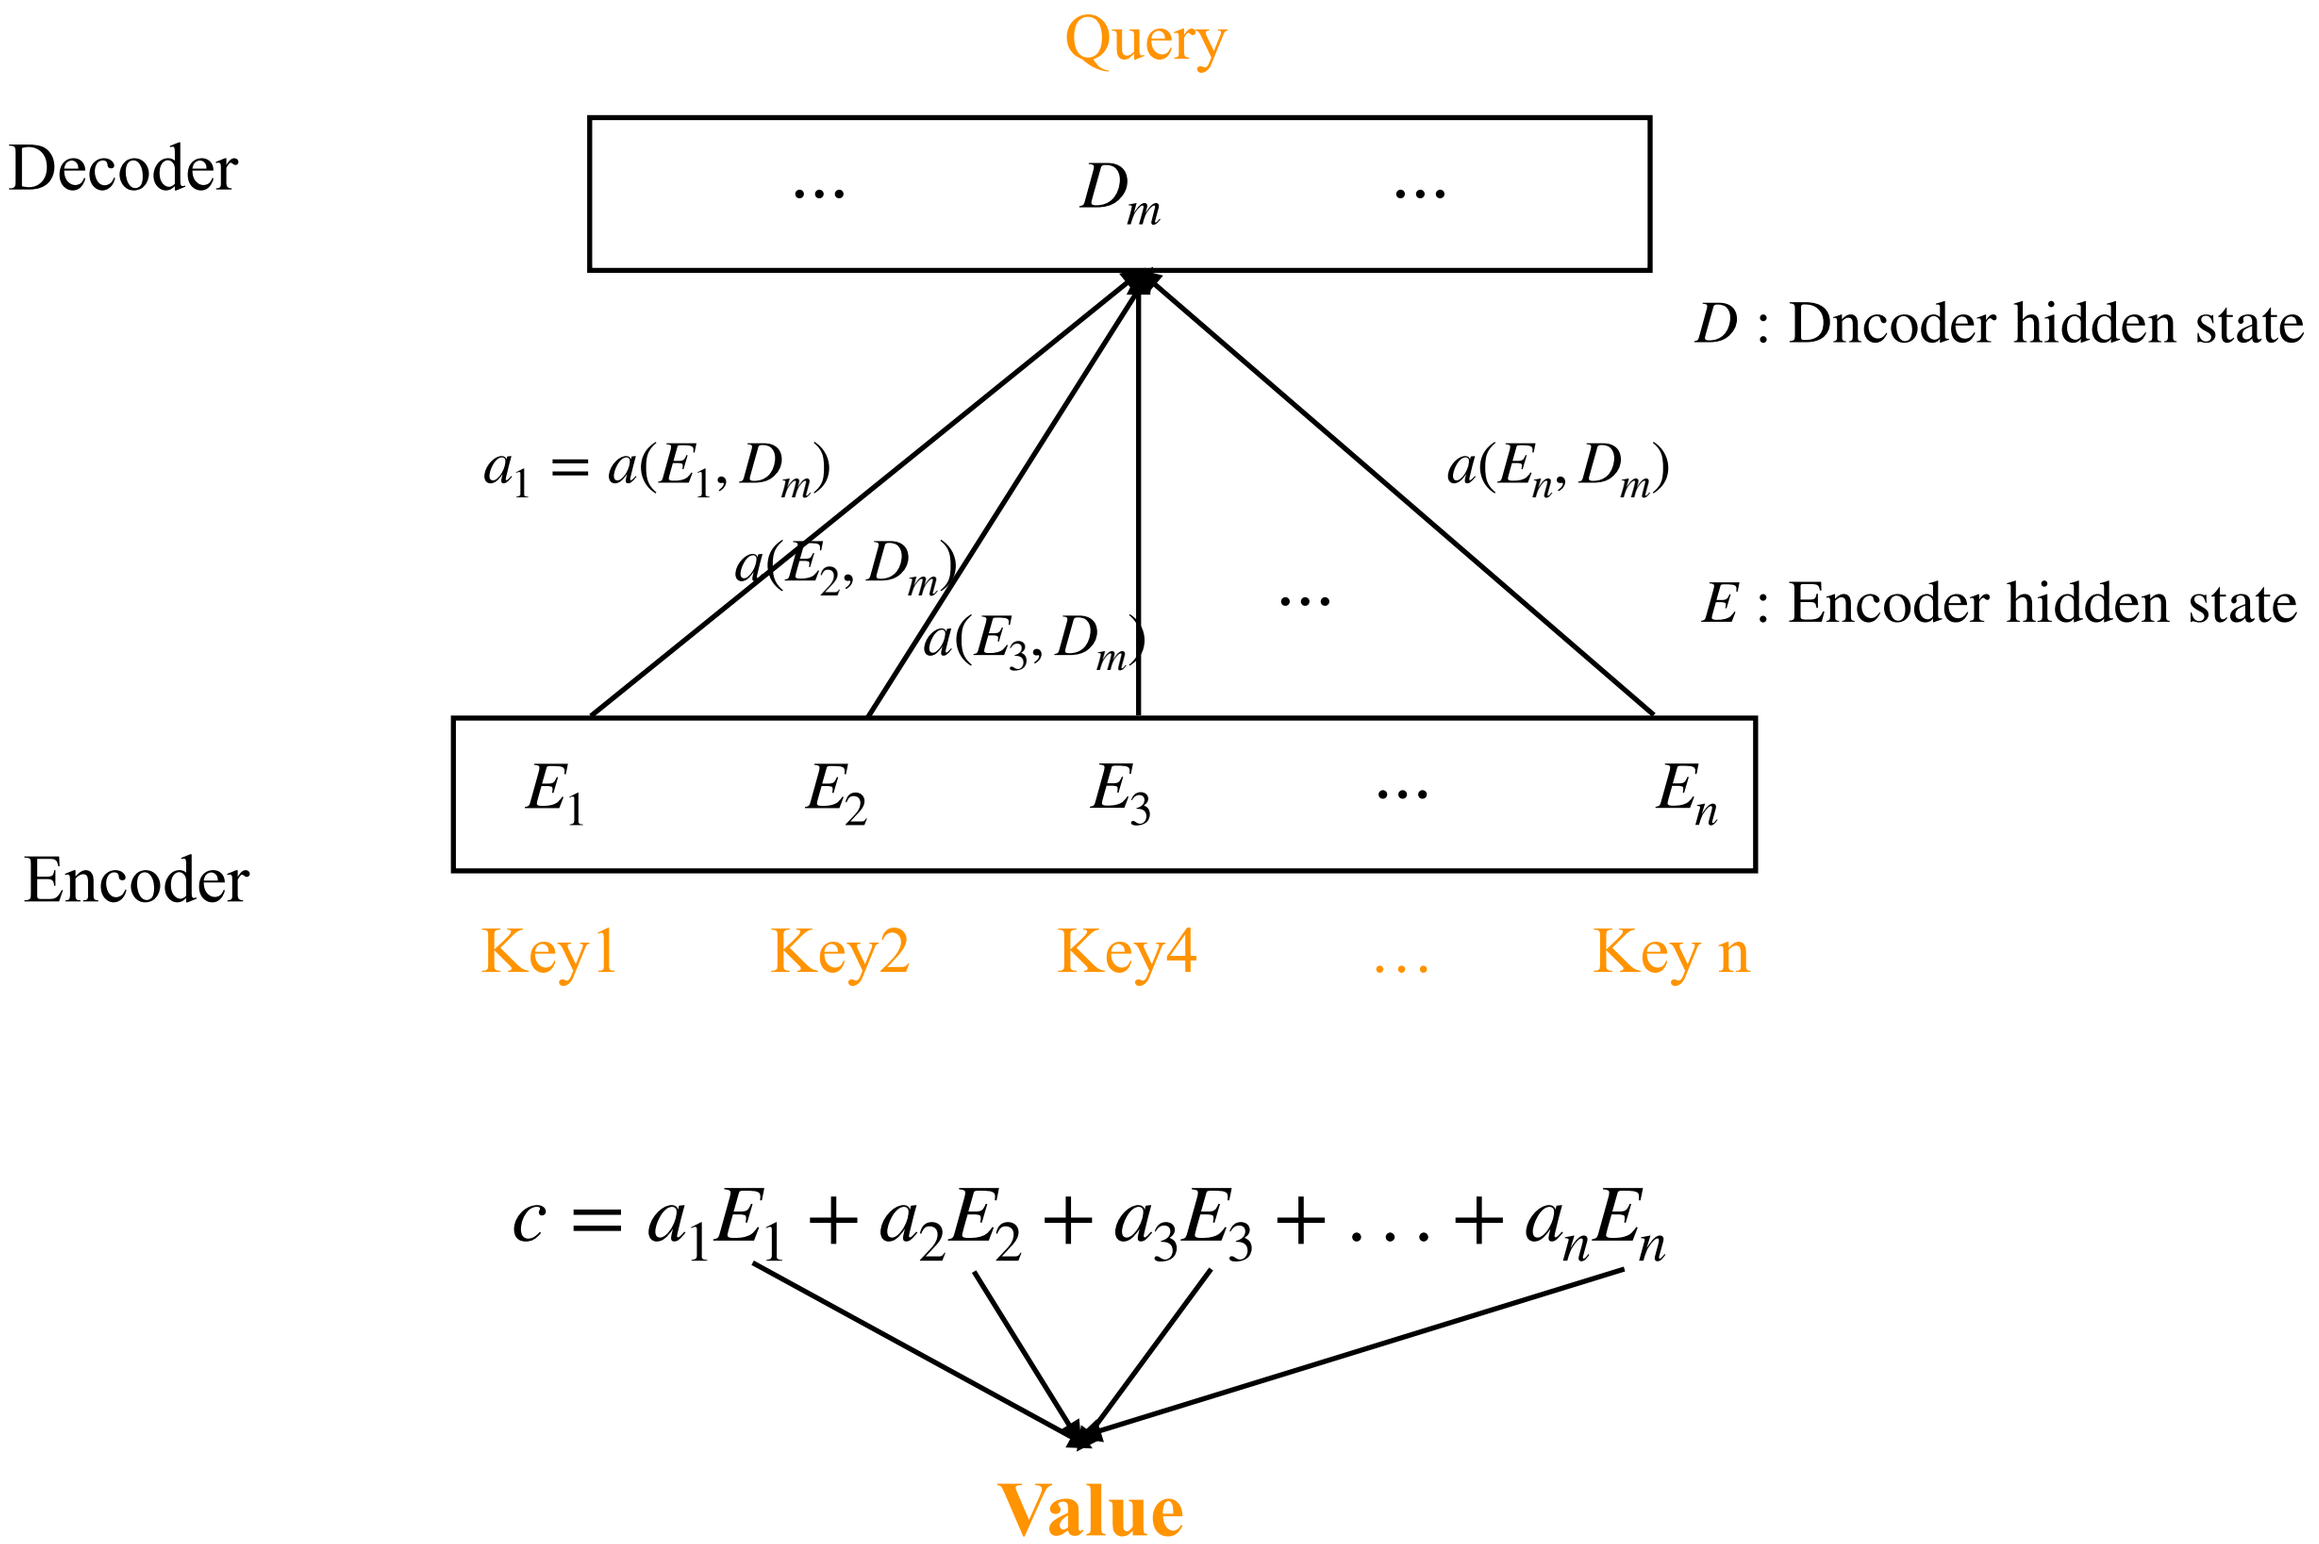

Transformer

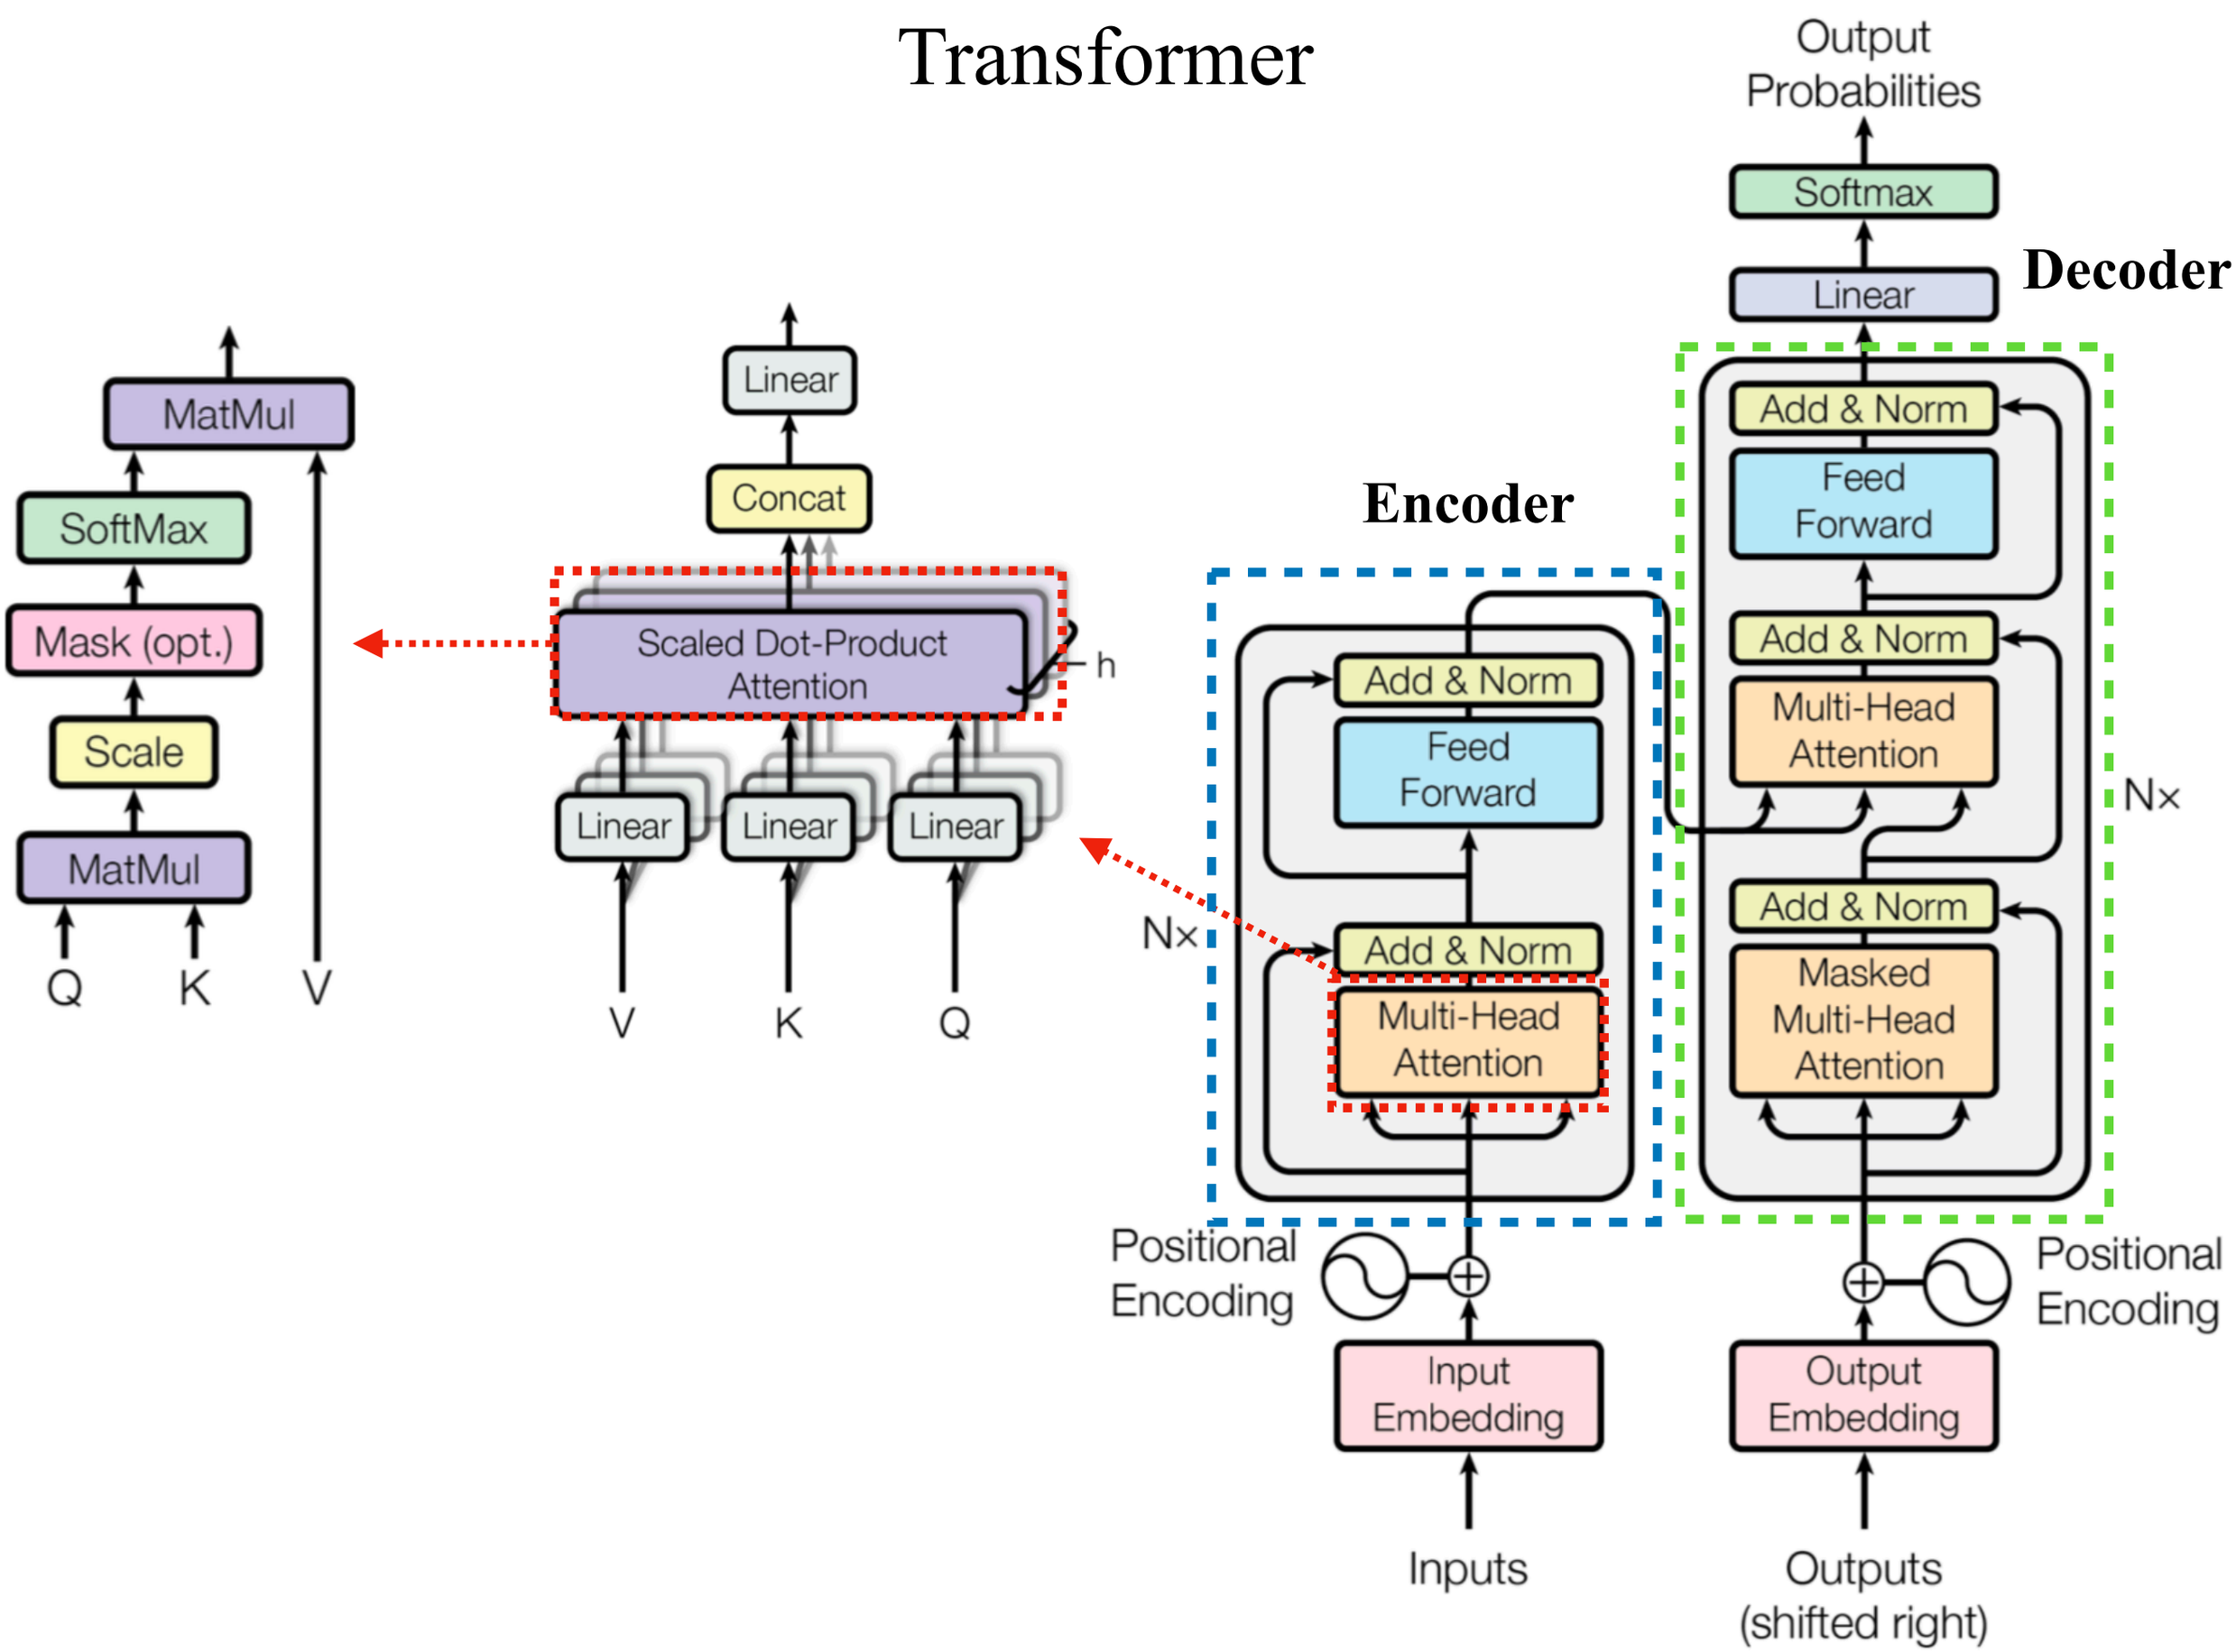

Transformer

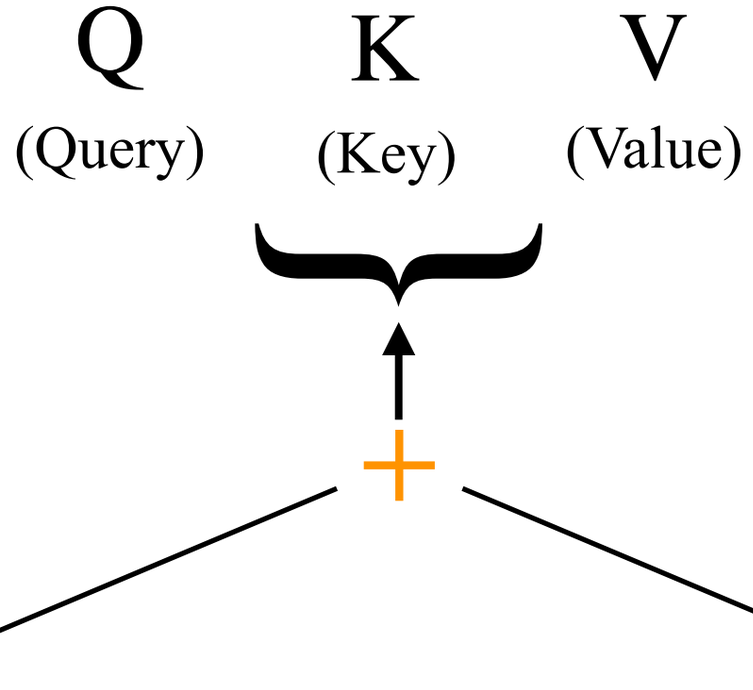

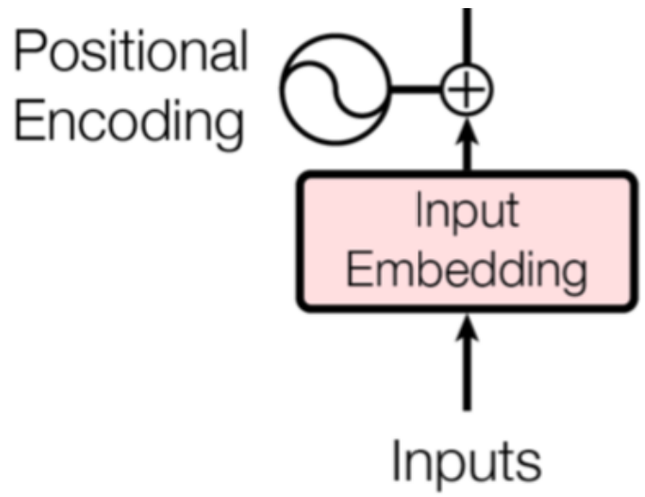

Positional Encoding

$$PE = \{p_0, p_1, p_2, \ldots, p_{n-1}\}$$

$$\begin{cases} PE_{(pos,2i)} = \sin(pos/10000^{\frac{2i}{d_{model}}}) \\ PE_{(pos,2i+1)} = \cos(pos/10000^{\frac{2i}{d_{model}}}) \end{cases}$$

Word Embedding

$$W = \{w_1, w_2, \ldots, w_n\},\ w_i \in d_{model}$$

Positional Encoding

Positional Encoding: $\{p_0, p_1, \ldots, p_{49}\}$, $p_i \in d_{model}$

1. Calculates the inner product p_{24} to the others p_0 *to* p_{49}.
2. Get product $V_0, V_1, \ldots, V_{49}$.

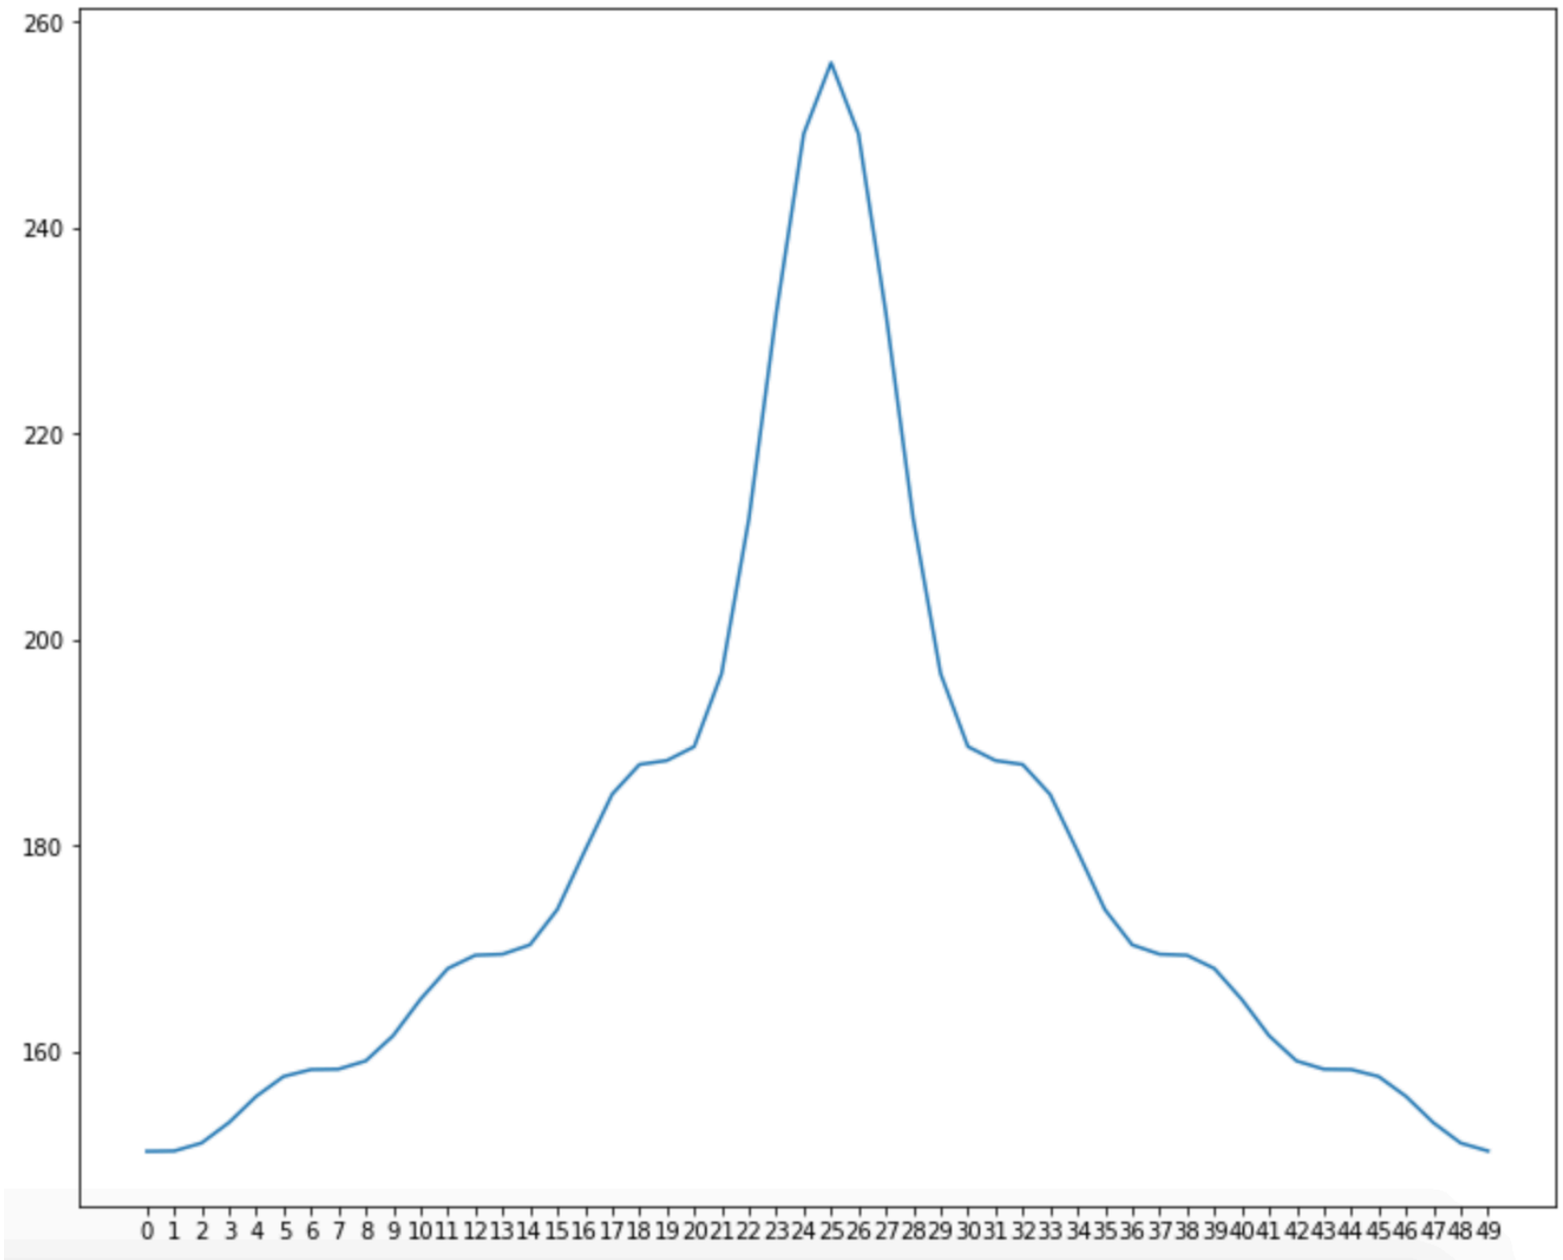

V

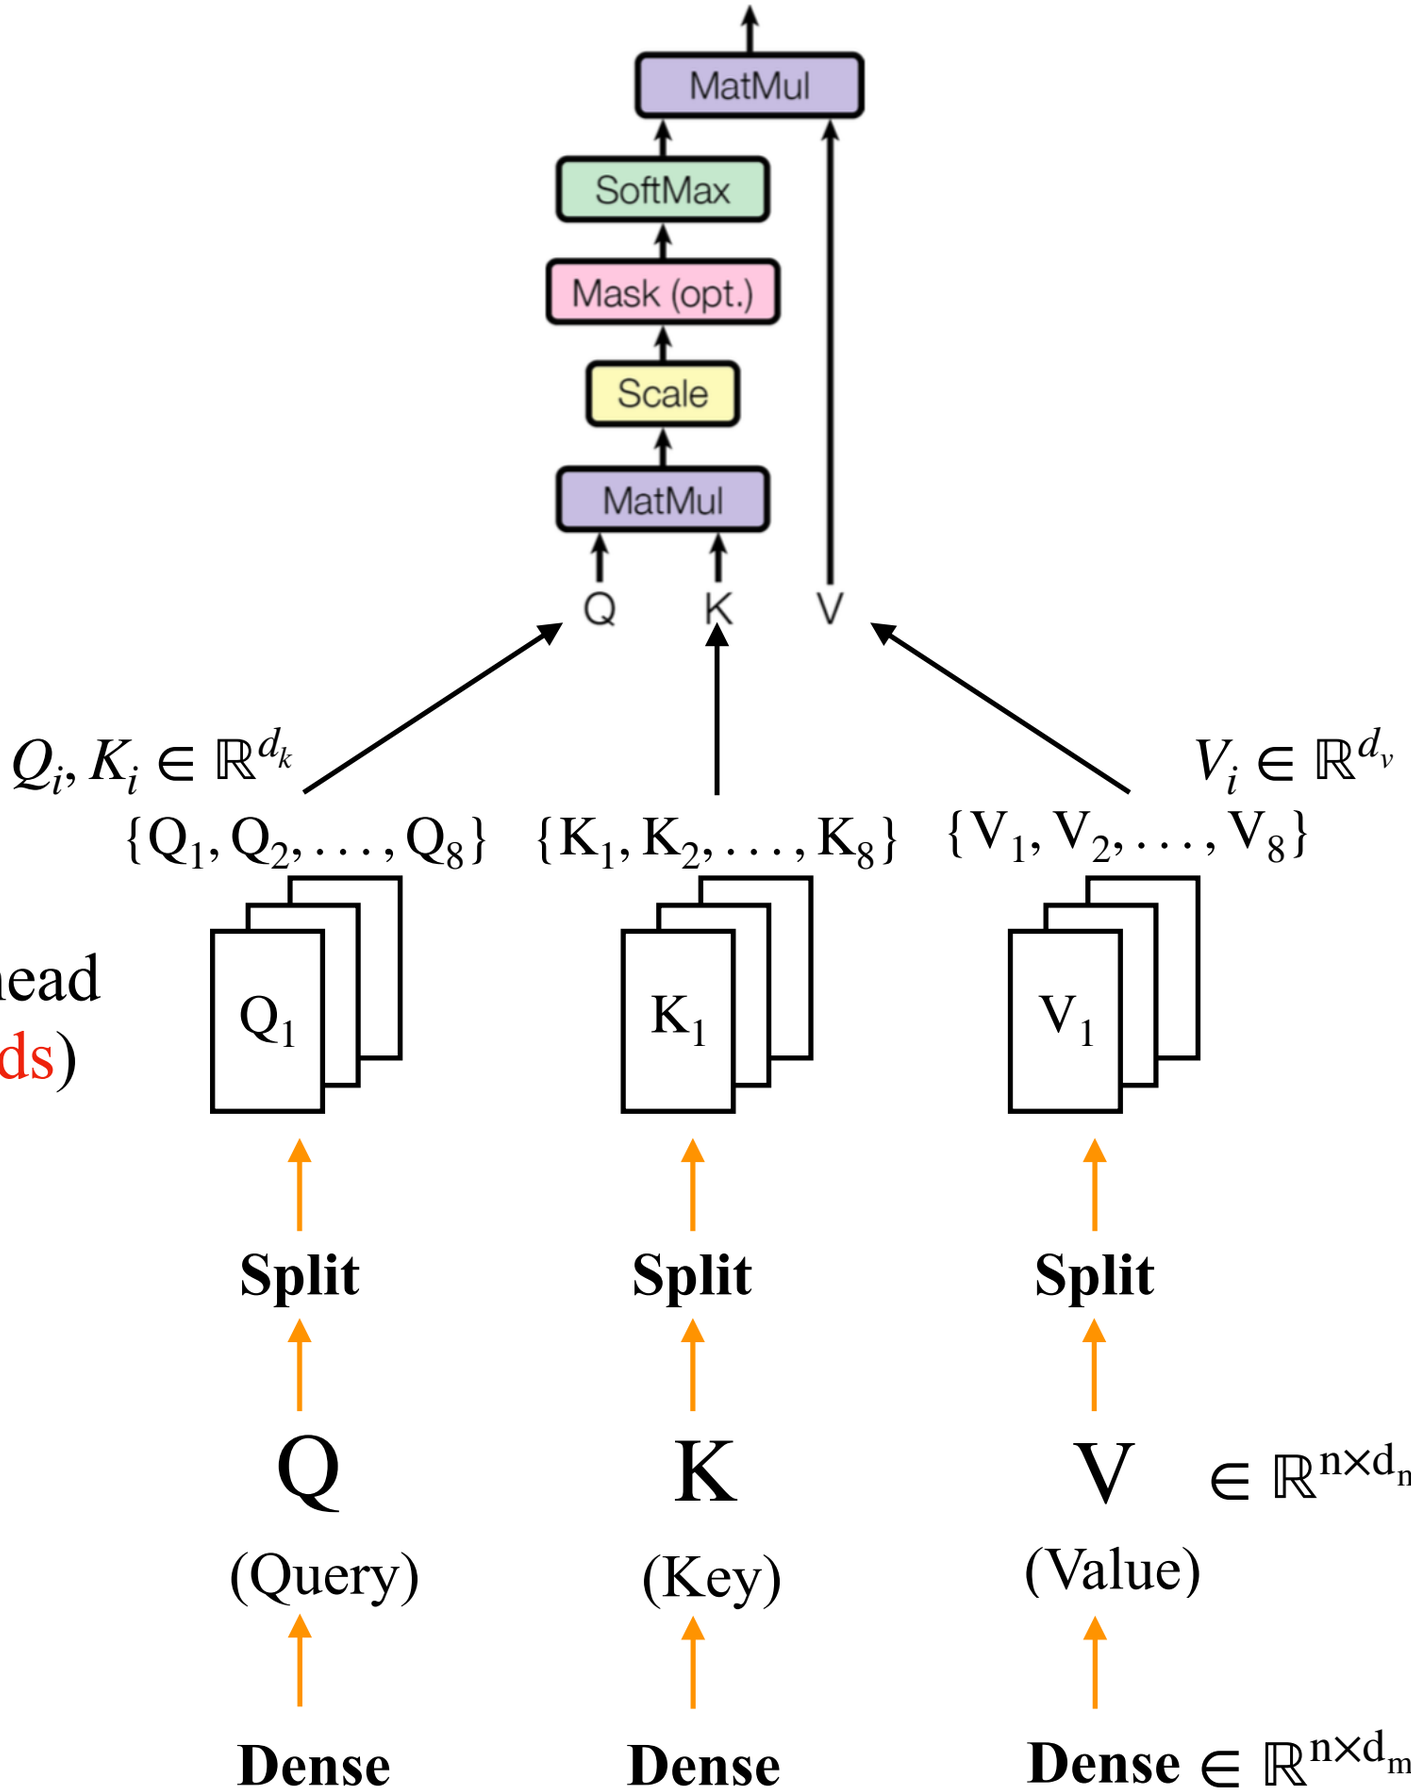
MatMul
SoftMax
Mask (opt.)
Scale
MatMul
Q
K
V
$Q_i, K_i \in \mathbb{R}^{d_k}$
$V_i \in \mathbb{R}^{d_v}$
$\{Q_1, Q_2, \ldots, Q_8\}$
$\{K_1, K_2, \ldots, K_8\}$
$\{V_1, V_2, \ldots, V_8\}$
Multi-head
(8-heads)
Q_1
K_1
V_1
Split
Split
Split
Q
K
V $\in \mathbb{R}^{n \times d_{model}}$
(Query)
(Key)
(Value)
Dense
Dense
Dense $\in \mathbb{R}^{n \times d_{model}}$

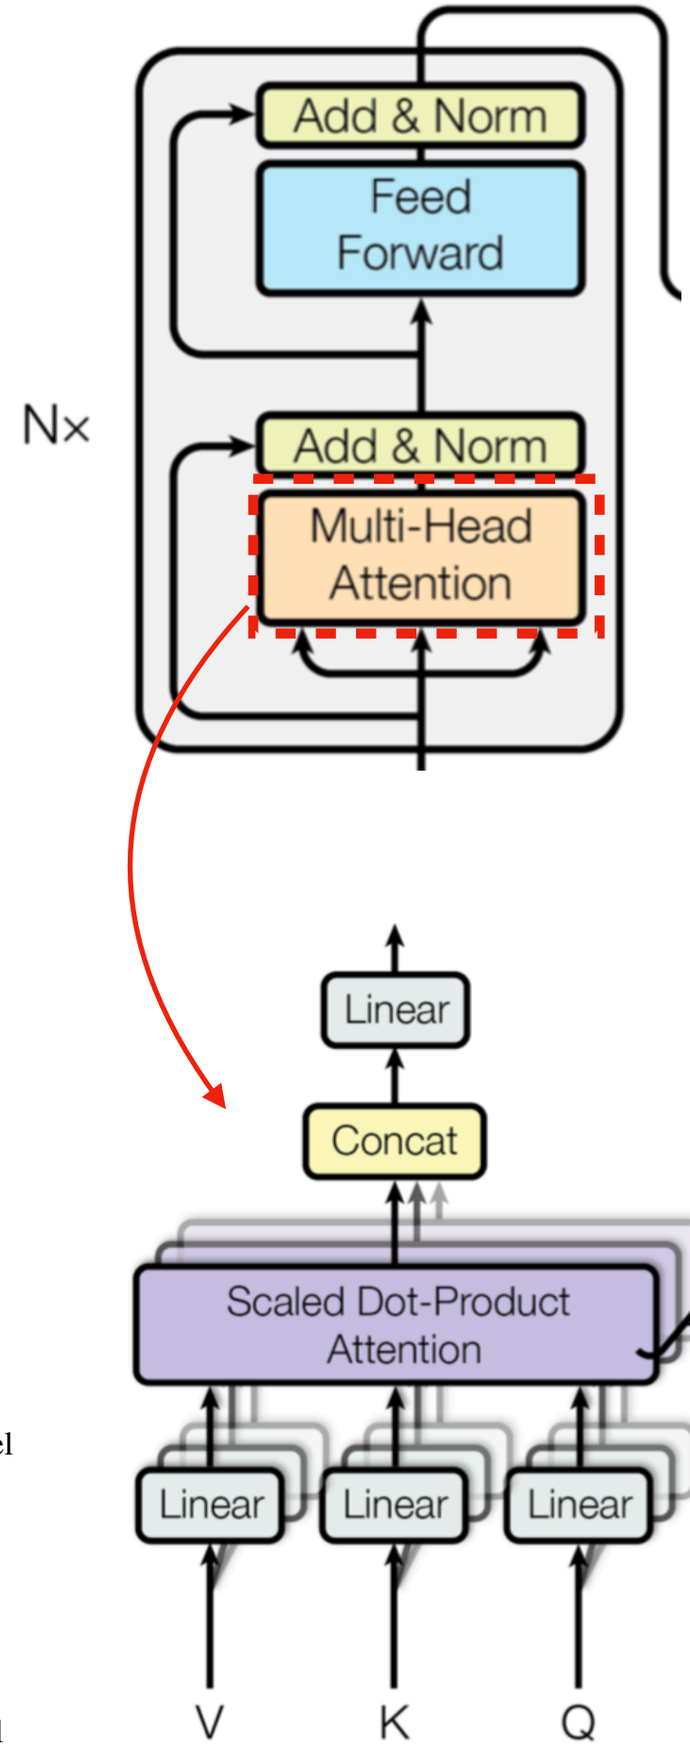
Nx
Add & Norm
Feed
Forward
Add & Norm
Multi-Head
Attention
Linear
Concat
Scaled Dot-Product
Attention
h
Linear
Linear
Linear
V
K
Q

Transformer

$$\text{Attention}(Q, K, V) = \text{softmax}(\frac{QK^T}{\sqrt{d_k}})V$$

$\in \mathbb{R}^{n \times n}$ • $\in \mathbb{R}^{n \times d_v}$ → $\in \mathbb{R}^{n \times d_v}$

Softmax

Scale Divided by $\sqrt{d_k}$

MatMul $\in \mathbb{R}^{n \times n}$ *Self-attention*

One-head $\{Q_1\}$ $\{K_1\}$ $\{V_1\}$

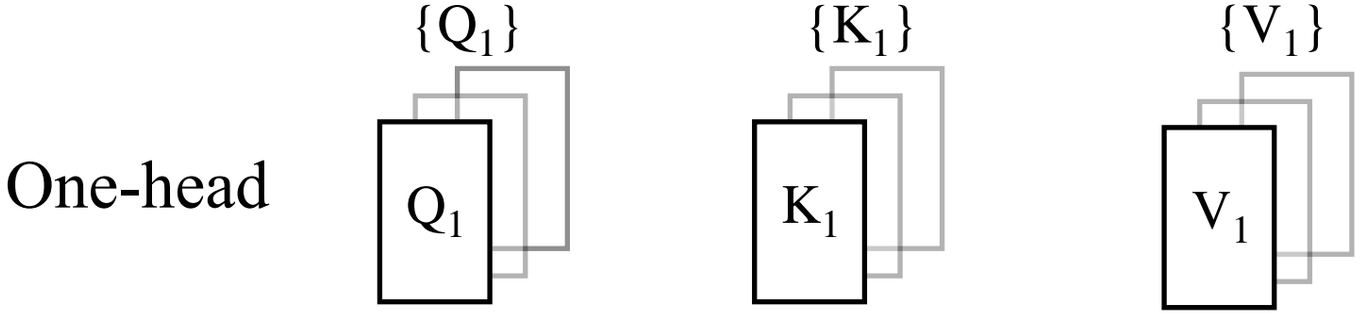

Multi-heads attention

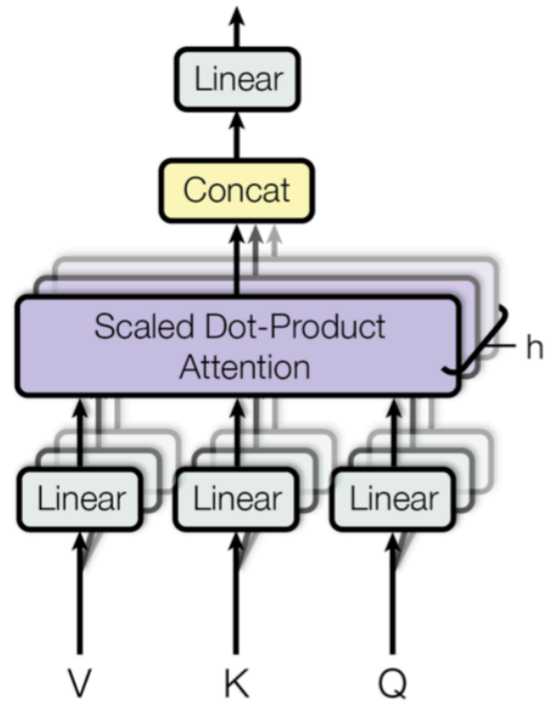

Scaled Dot-Product Attention (Self-Attention)

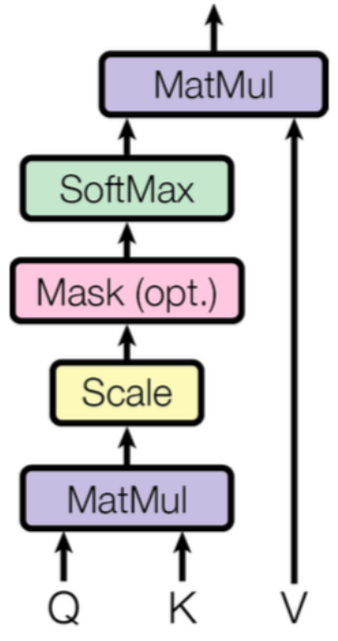

Transformer

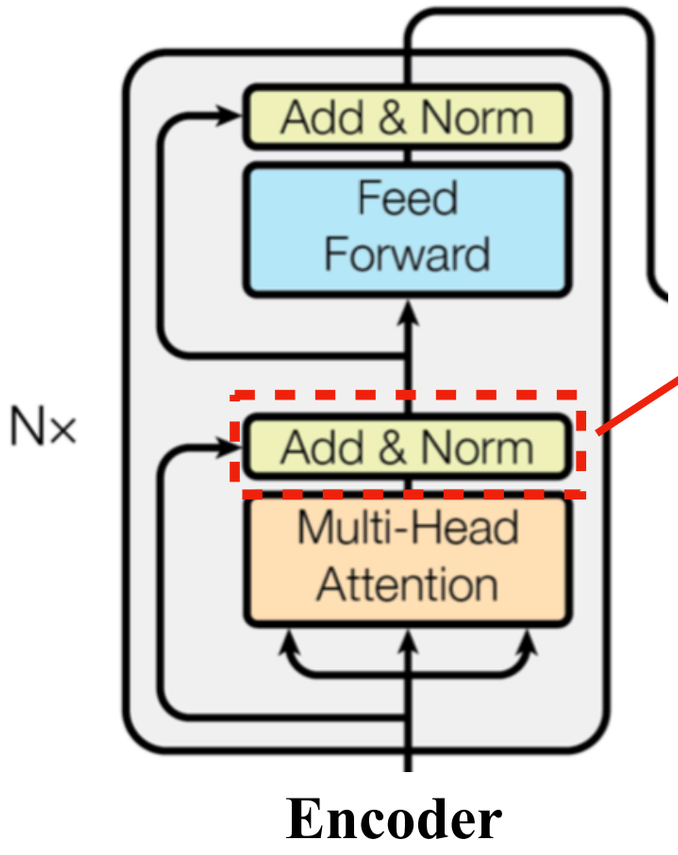

Encoder

Residual connection

$$\text{Normalize}\Big(\underbrace{\square\square \cdots \square}_{\in \mathbb{R}^{n\times d_{model}}} + \underbrace{\text{Q (Query)}}_{\in \mathbb{R}^{n\times d_{model}}}\Big)$$

$$= \text{Normalize}\Big(\underbrace{\qquad\qquad}_{\in \mathbb{R}^{n\times d_{model}}}\Big)$$

$$\text{Layer Normalize: } LN(z;\alpha,\beta) = \frac{(z-\mu)}{\sigma} \odot \alpha + \beta$$

$$\text{Mean} : \mu^l = \frac{1}{D}\sum_{i=1}^{D} z_i^l \qquad \text{Gains} : \alpha$$

$$\text{Standard Deviation} : \sigma^l = \sqrt{\frac{1}{D}\sum_{i=1}^{D}(z_i^l - \mu^l)^2} \qquad \text{Biases} : \beta$$

Transformer

Feed Forward
(Dense、Conv1d)

$$\mathrm{FFN}(x) = max(0, xW_1 + b_1)W_2 + b_2$$

$$\mathrm{Dense}(d_{model}) \rightarrow \mathrm{Dense}\ (4 \times d_{model}) \rightarrow \mathrm{Dense}(d_{model})$$

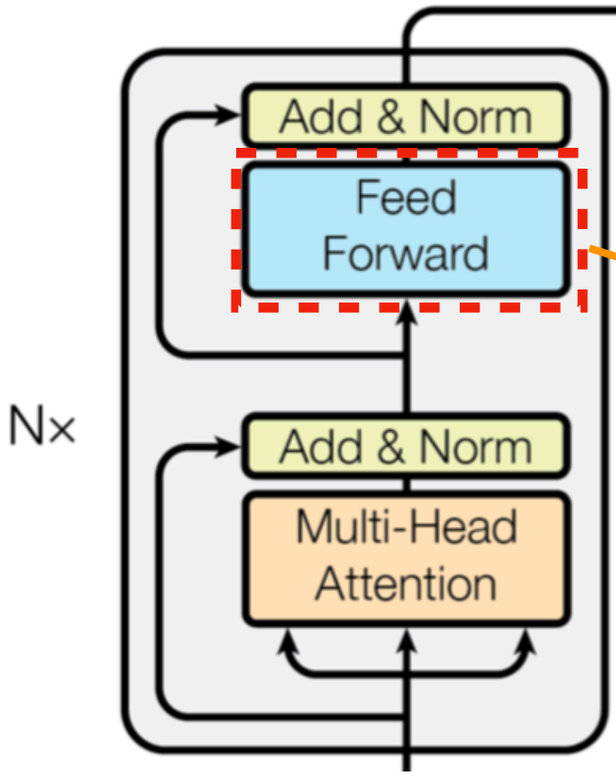

Encoder

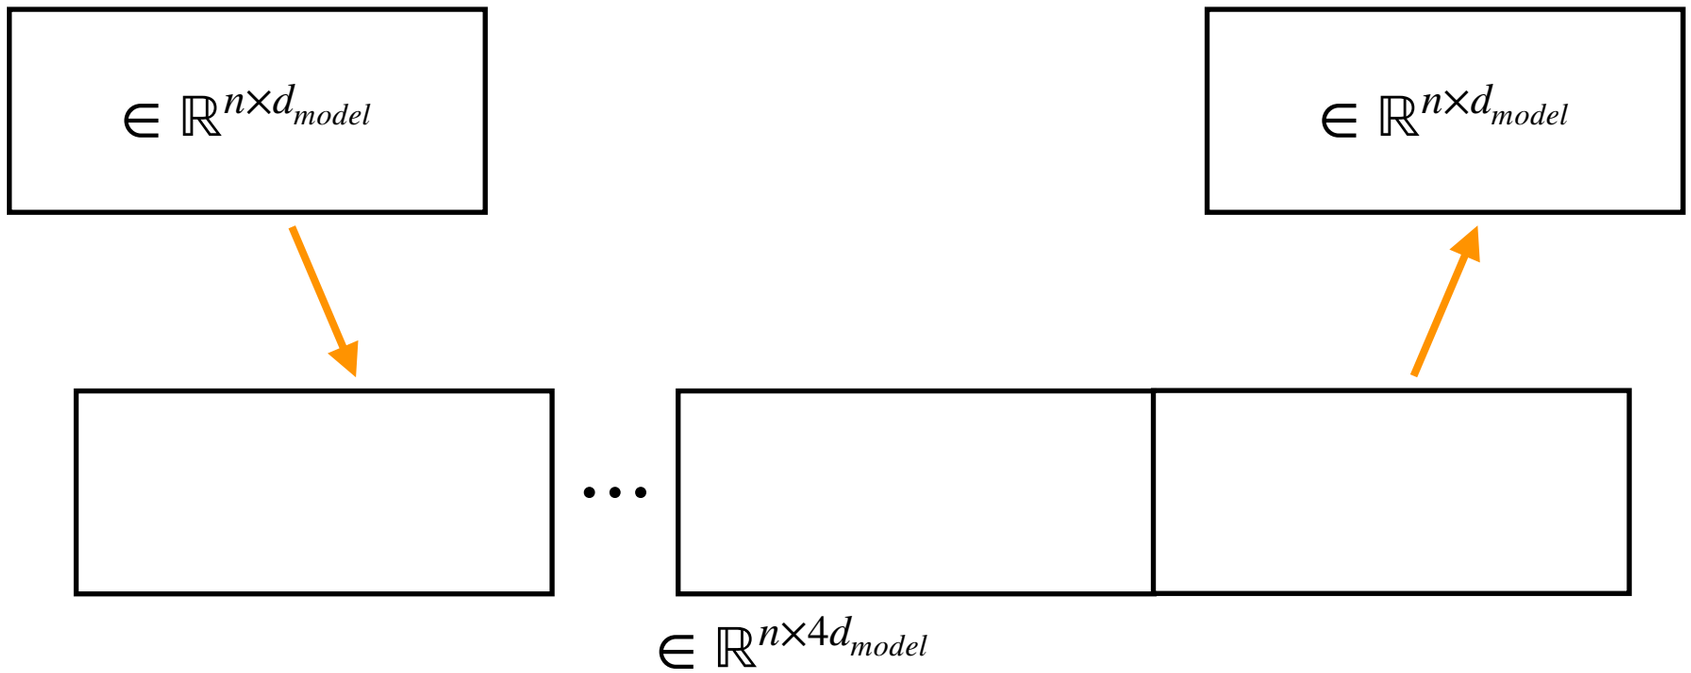

Transformer

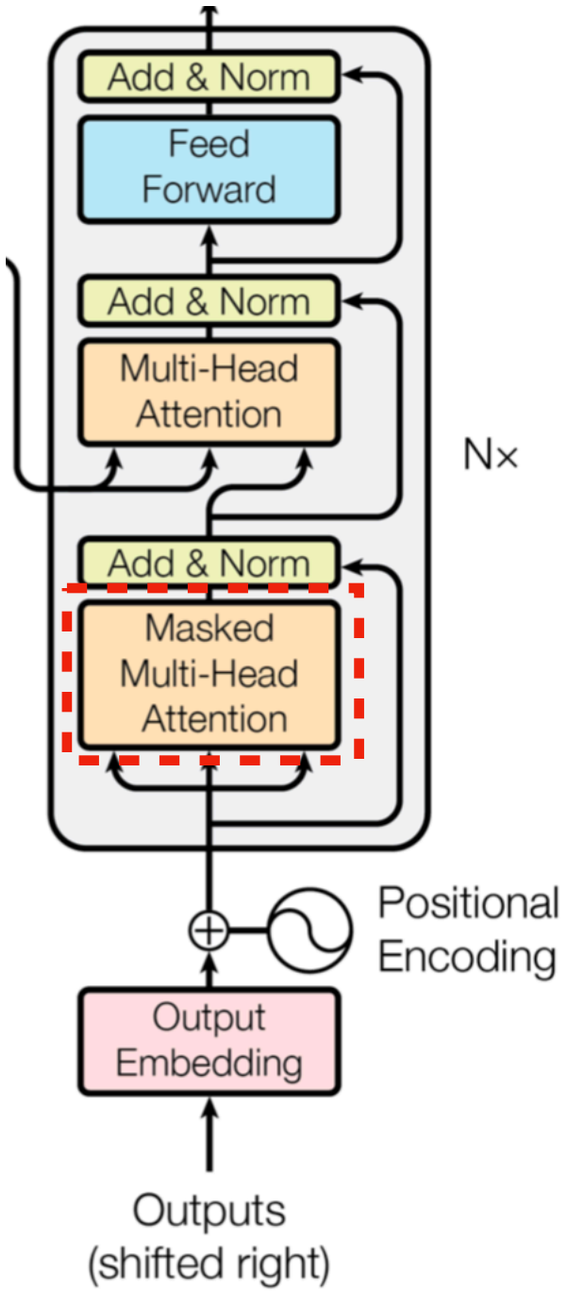

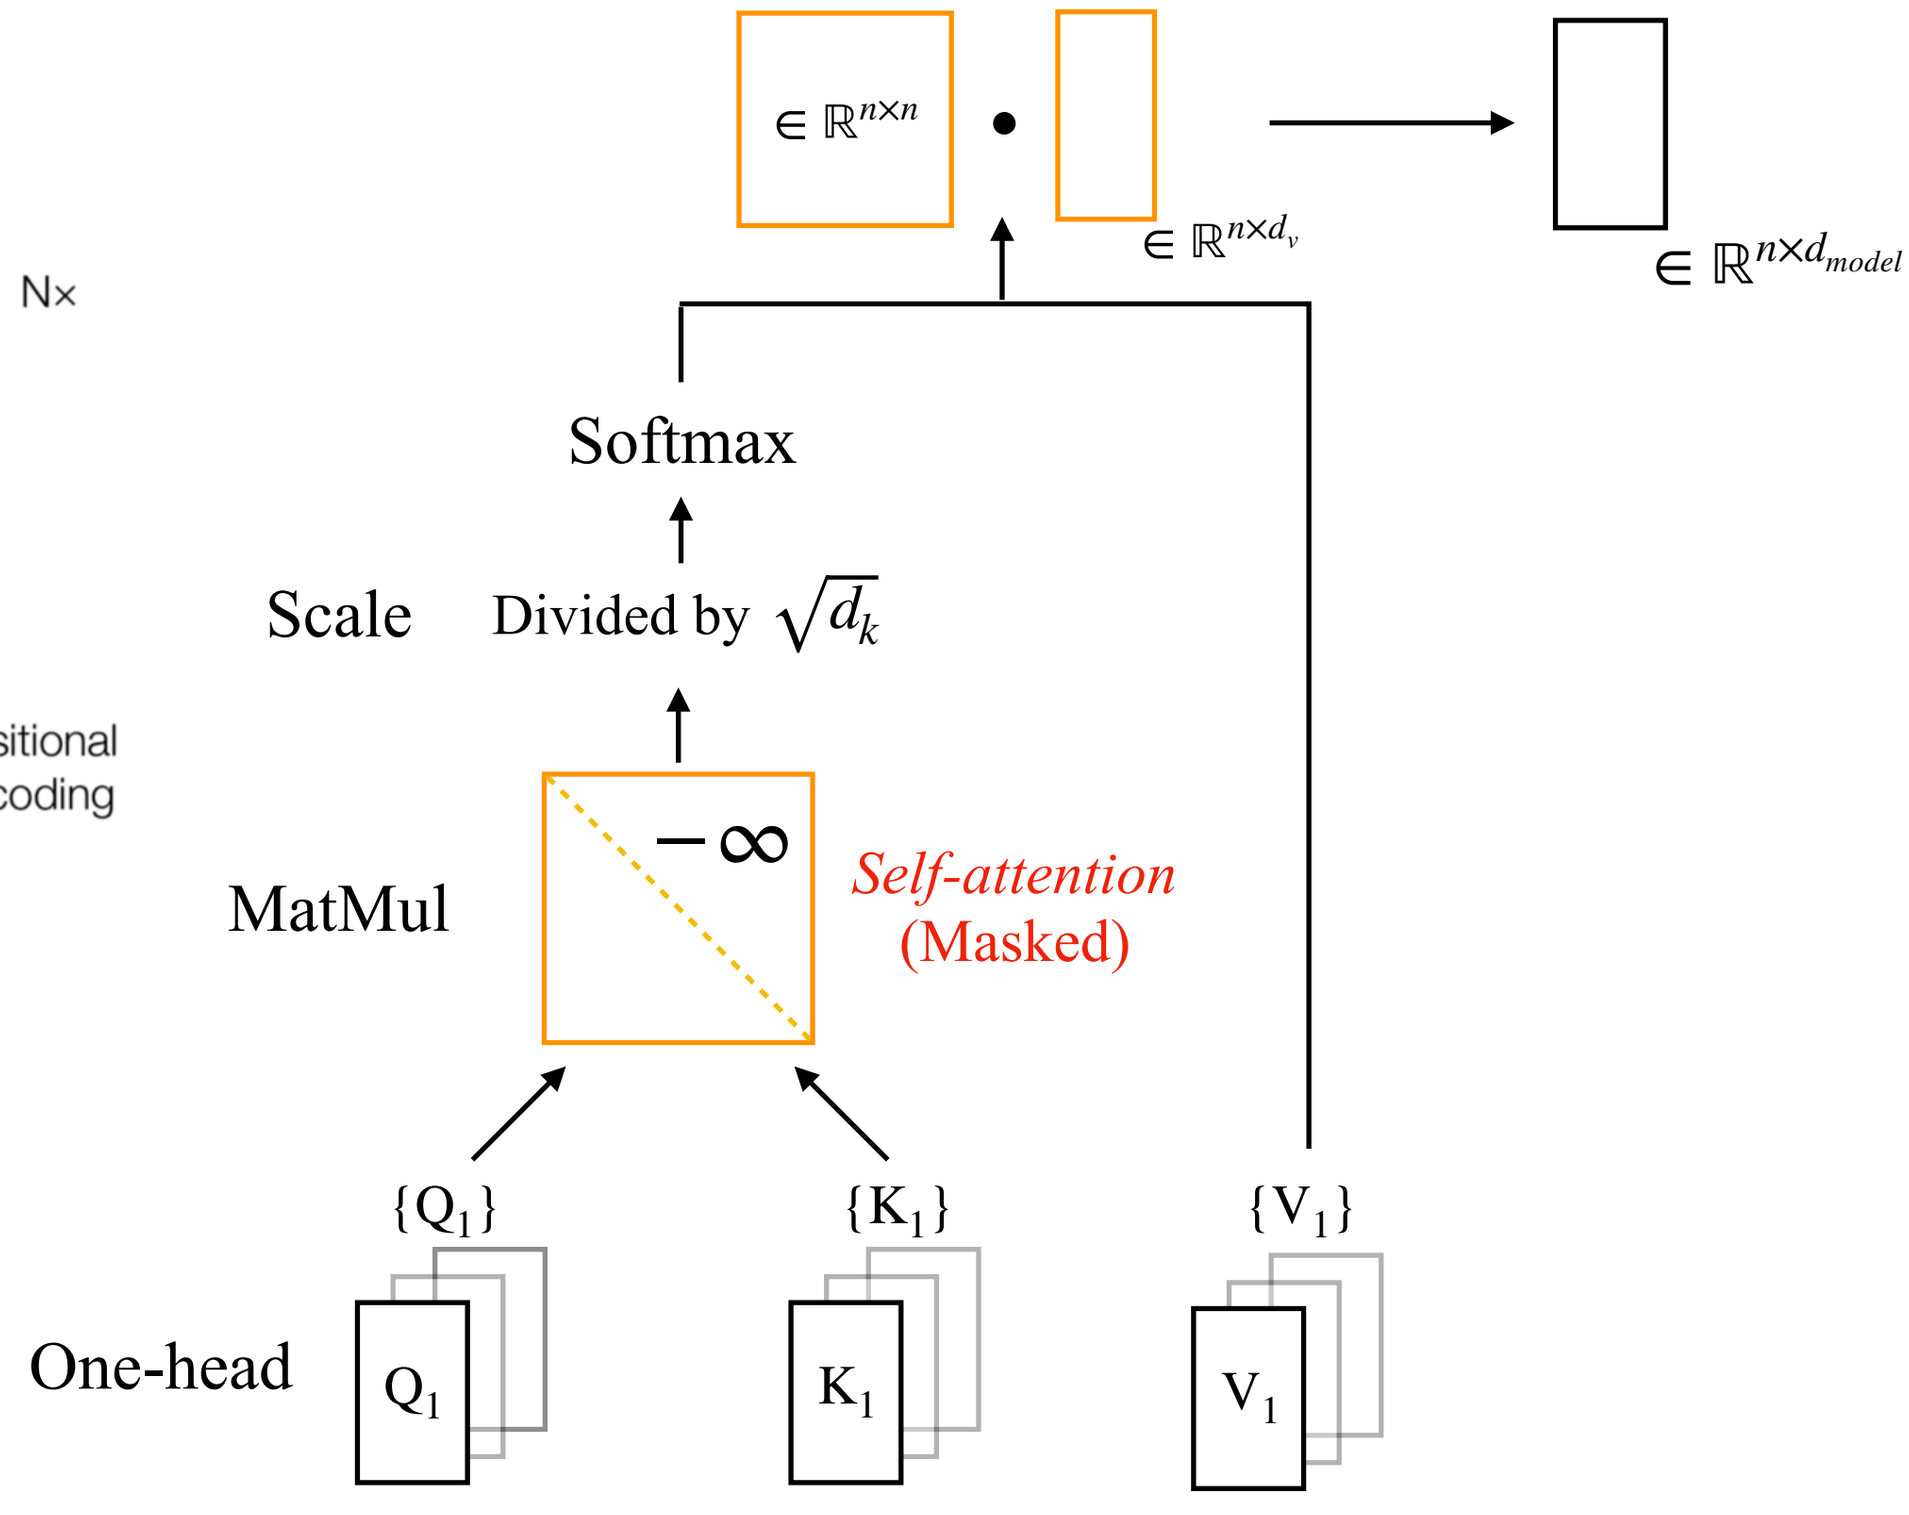

Transformer

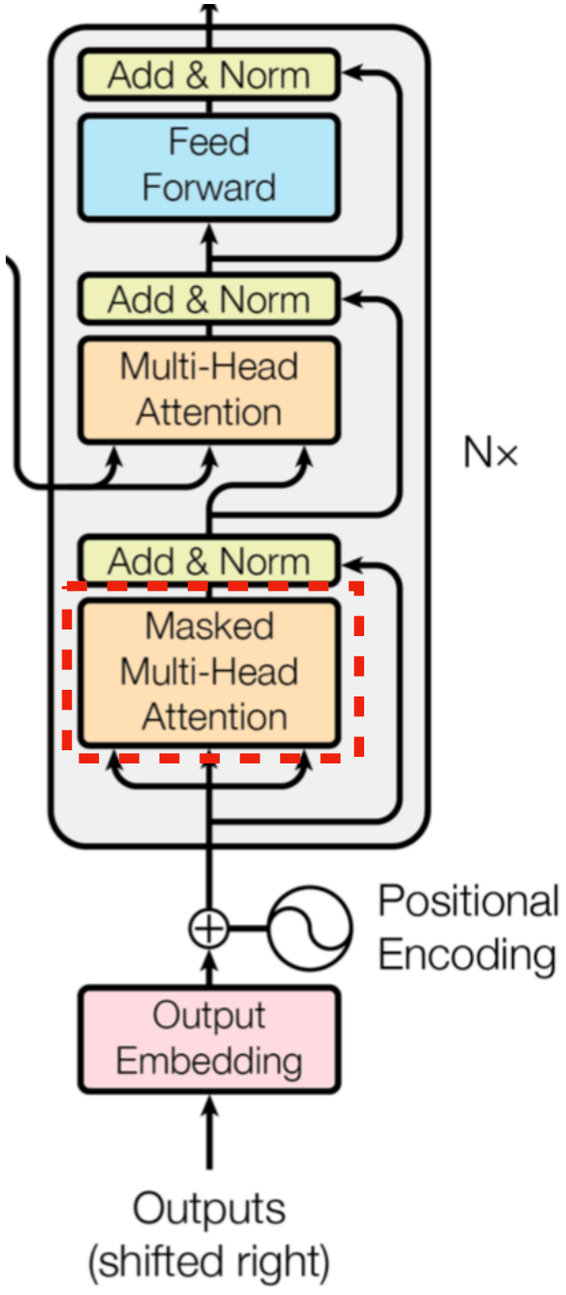

Decoder

Softmax → 0

↑

Scale Divided by $\sqrt{d_k}$

↑

We need to prevent leftward information flow in the decoder

$M_{1,1}$: The similarity between Q_1 and K_1

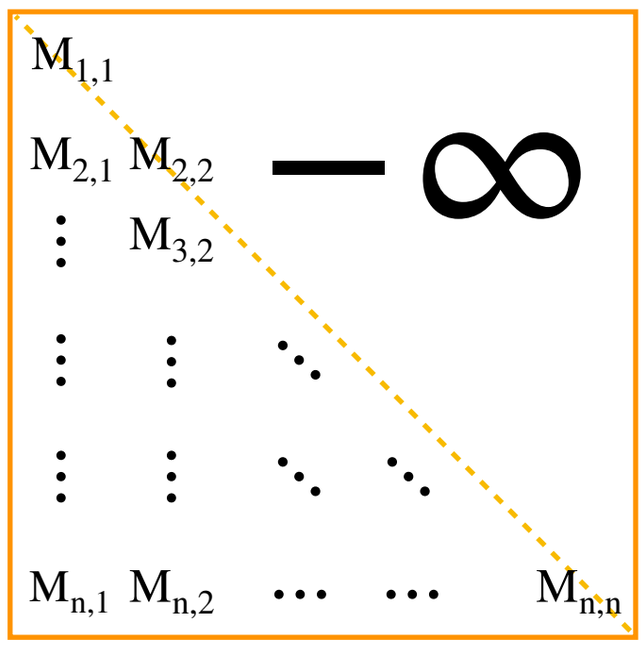

Self-attention (Masked)

w_1 w_2 w_3 w_4 w_5 . . . w_n

w_1 w_2 w_3 w_4 w_5 . . . w_n

w_1 w_2 w_3 w_4 w_5 . . . w_n

⋮

w_1 w_2 w_3 w_4 w_5 . . . w_n

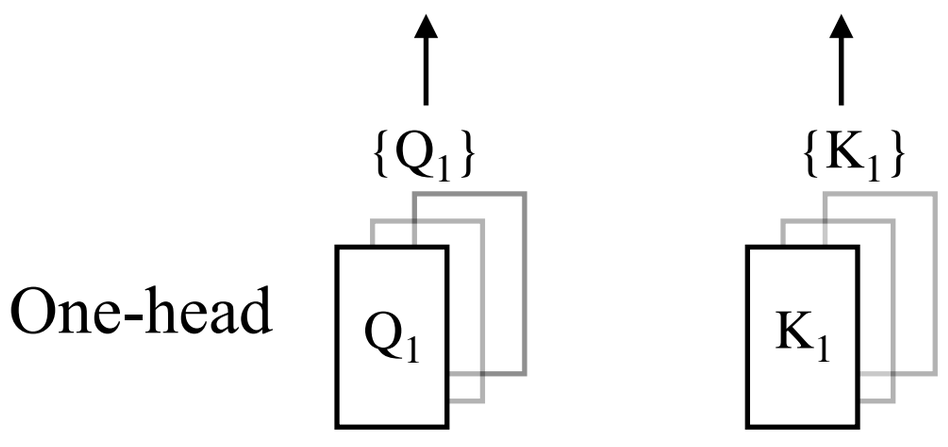

Evaluation

Table 2: The Transformer achieves better BLEU scores than previous state-of-the-art models on the English-to-German and English-to-French newstest2014 tests at a fraction of the training cost.

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	**41.29**	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	**$3.3 \cdot 10^{18}$**	
Transformer (big)	**28.4**	**41.8**	$2.3 \cdot 10^{19}$	

- **ByteNet:** 2 convolution layers
- **Deep-Att + PosUnk:** 2 Bi-LSTM layers(Encoder) + 1 LSTM layer (Decoder)
- **GNMT + RL:** 7 LSTM layers + 1 Bi-LSTM layer(Encoder) + 8 LSTM layers(Decoder)
- **Transformer(base):** training 100,000 steps(12 hours), 0.4 seconds per steps
- **Transformer(big):** training 300,000 steps(3.5 days), 1.0 seconds per steps

Evaluation

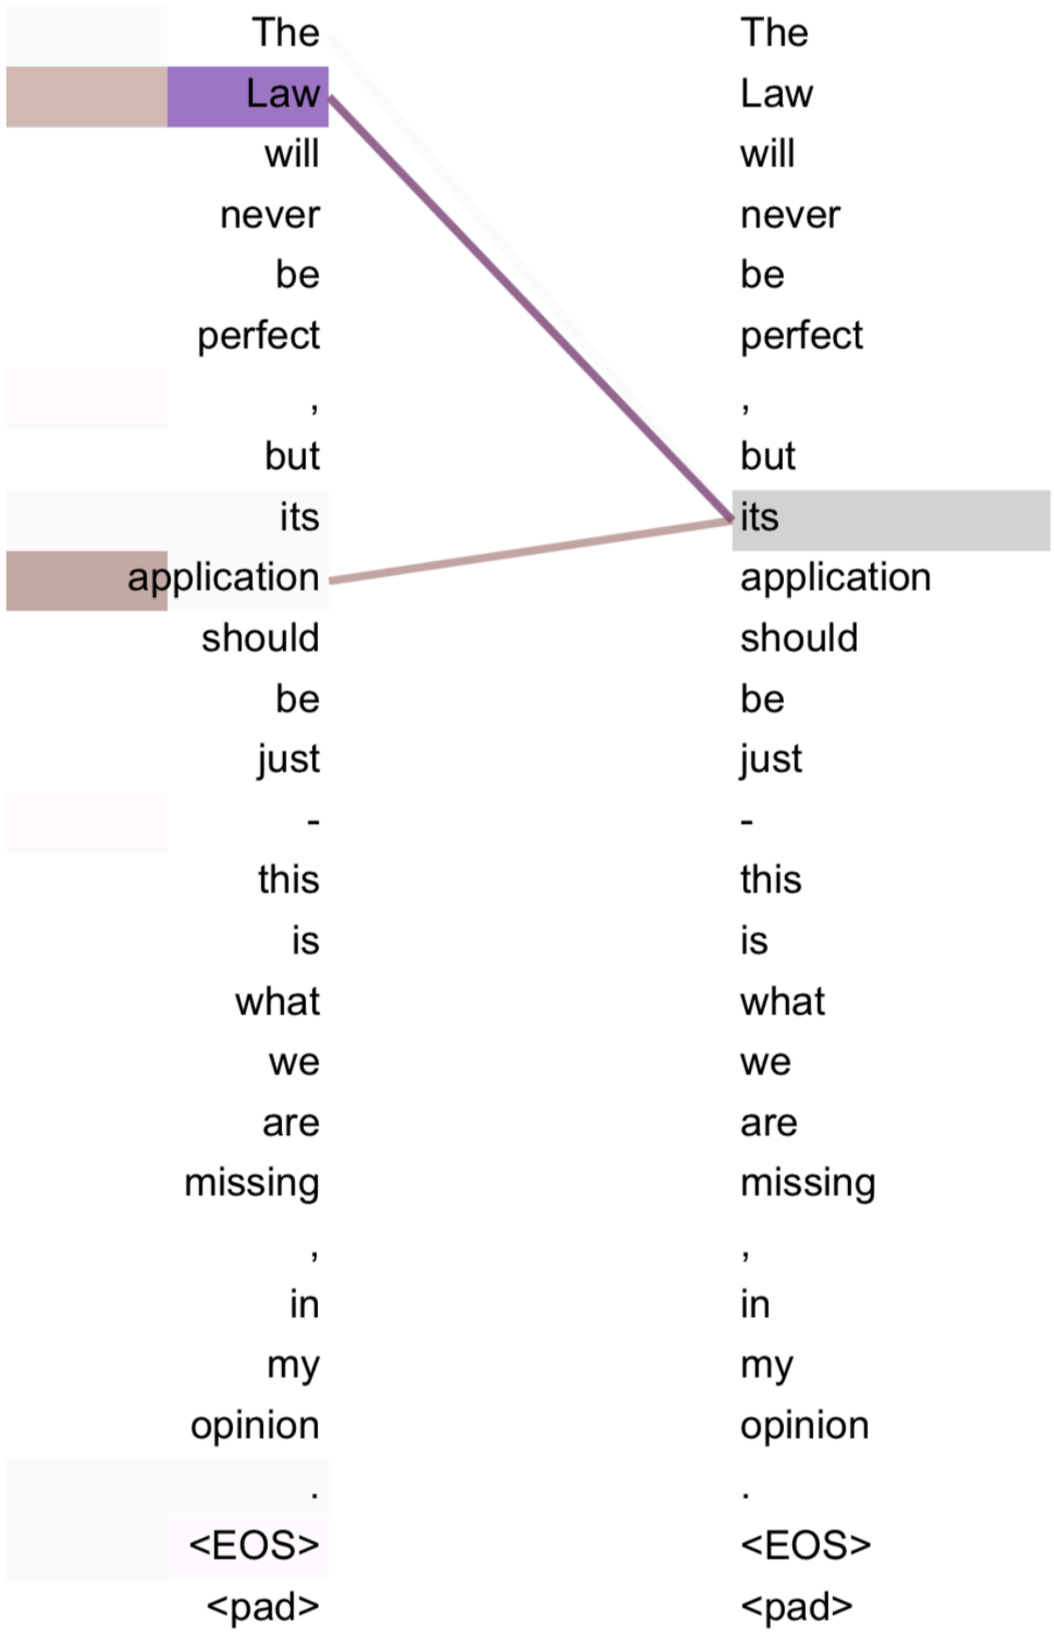
The
Law
will
never
be
perfect
,
but
its
application
should
be
just
-
this
is
what
we
are
missing
,
in
my
opinion
.
<EOS>
<pad>
The
Law
will
never
be
perfect
,
but
its
application
should
be
just
-
this
is
what
we
are
missing
,
in
my
opinion
.
<EOS>
<pad>

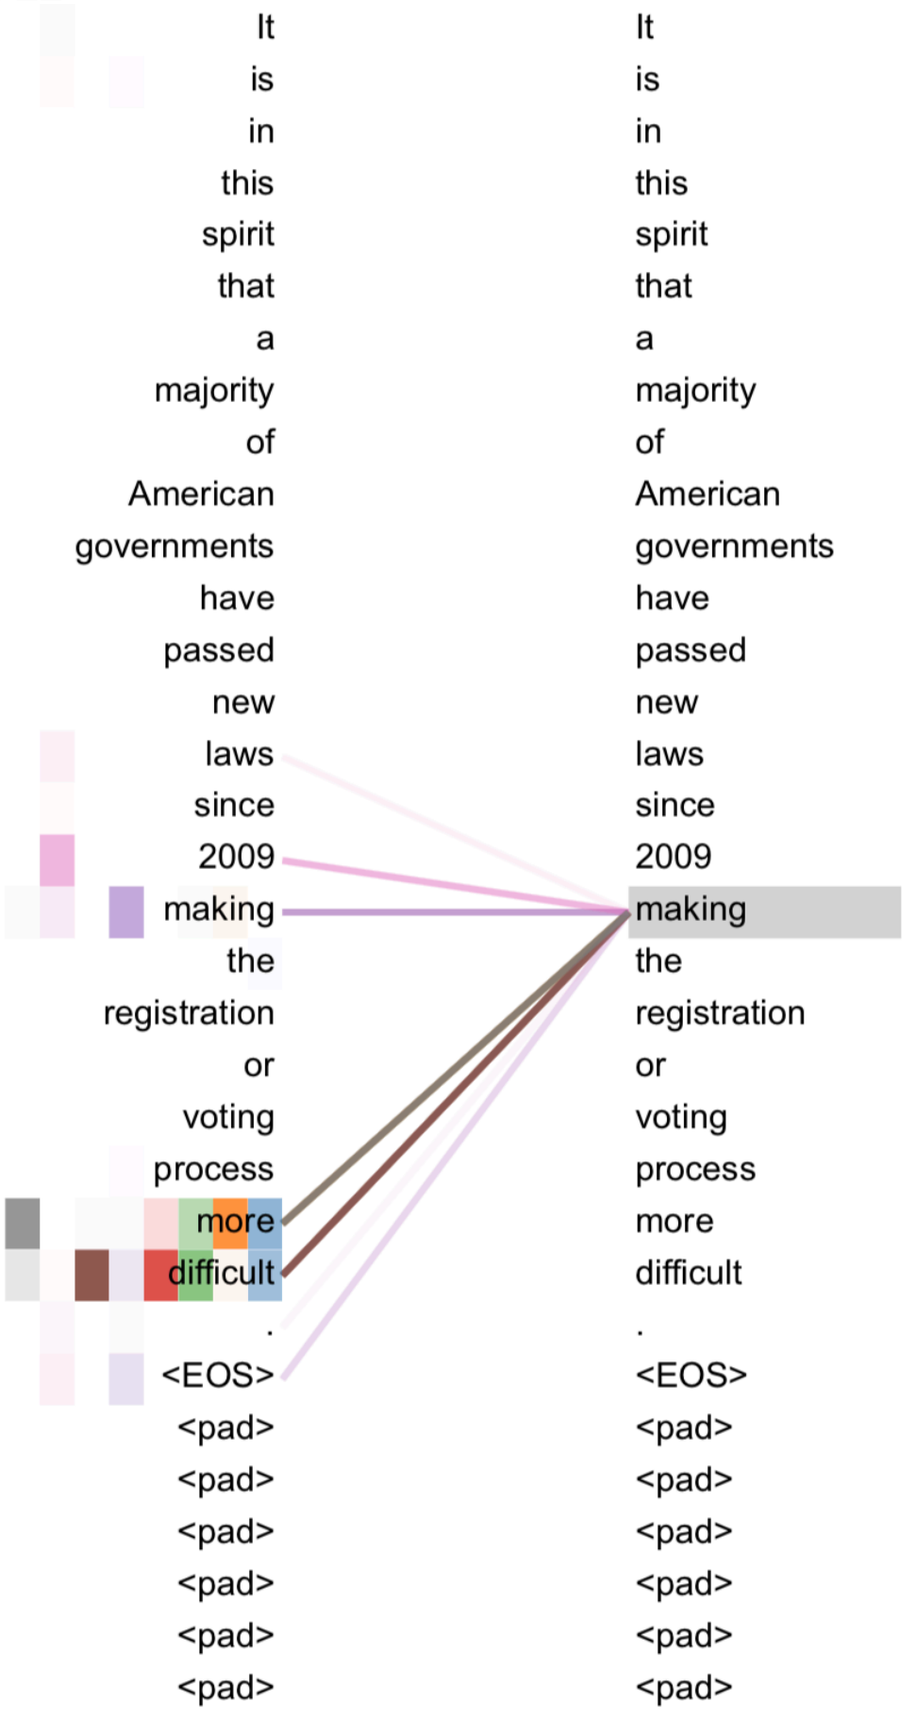
It
is
in
this
spirit
that
a
majority
of
American
governments
have
passed
new
laws
since
2009
making
the
registration
or
voting
process
more
difficult
.
<EOS>
<pad>
<pad>
<pad>
<pad>
<pad>
<pad>
It
is
in
this
spirit
that
a
majority
of
American
governments
have
passed
new
laws
since
2009
making
the
registration
or
voting
process
more
difficult
.
<EOS>
<pad>
<pad>
<pad>
<pad>
<pad>
<pad>

Evaluation

Source sentence:
Taiwan is a beautiful country.

Predict sentence:
台湾是一个美丽的国家。